Q1 2026 Earnings Call

Company Participants

- Amy Hood, Executive Vice President & Chief Financial Officer
- Jonathan Neilson, Vice President, Investor Relations
- Satya Nadella, Chairman & Chief Executive Officer

Other Participants

- Brad Zelnick, Analyst, Deutsche Bank
- Brent Thill, Analyst, Jefferies
- Karl Keirstead, Analyst, UBS
- Kash Rangan, Analyst, Goldman Sachs
- Keith Weiss, Analyst, Morgan Stanley
- Mark Moerdler, Analyst, Bernstein
- Mark Murphy, Analyst, JP Morgan

Presentation

Operator

Greetings, and welcome to the Microsoft Fiscal Year 2026 First Quarter Earnings Conference Call. At this time, all participants are in a listen-only mode. A question-and-answer session will follow the formal presentation. (Operator Instructions). As a reminder, this conference is being recorded.

It is now my pleasure to introduce Jonathan Neilson, Vice President of Investor Relations. Please go ahead.

Jonathan Neilson

Good afternoon and thank you for joining us today. On the call with me are Satya Nadella, Chairman and Chief Executive Officer; Amy Hood, Chief Financial Officer; Alice Jolla, Chief Accounting Officer and Keith Dolliver, Corporate Secretary and Deputy General Counsel.

On the Microsoft Investor Relations website, we will provide our earnings press release and financial summary slide deck, which is intended to supplement our prepared remarks and provides the reconciliation of differences between GAAP and non-GAAP financial measures. More detailed outlook slides will be available on the Microsoft Investor Relations website.

On this call, we will discuss certain non-GAAP items. The non-GAAP financial measures provided should not be considered as a substitute for or superior to the measures of financial performance prepared in accordance with GAAP. They are included as additional clarifying items to aid investors in further understanding the company's first quarter performance in addition to the impact these items and events have on the financial results. All growth comparisons we make on the call today relates to the corresponding period of last year unless otherwise noted.

We will also provide growth rates in constant currency when available as a framework for assessing how our underlying business performed, excluding the effect of foreign currency rate fluctuations. Where growth rates are the same in constant currency, we will refer to the growth rate only. We will post our prepared remarks to our website.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

Today's call is being recorded. If you ask a question, it will be included in our live transmission, in the transcript and in any future use of the recording. You can replay the call and view the transcript on the Microsoft Investor Relations website.

During this call, we will be making forward-looking statements, which are predictions, projections or other statements about future events. These statements are based on current expectations and assumptions that are subject to risks and uncertainties. Actual results could materially differ because of factors discussed in today's earnings press release, in the comments made during this conference call and in the Risk Factors section of our Form 10-K, Forms 10-Q and other reports and filings with the Securities and Exchange Commission. We do not undertake any duty to update any forward-looking statements.

And with that, I'll turn the call over to Satya.

Satya Nadella {BIO 3224315 <GO>}

Thank you, Jonathan. It is a very strong start to our fiscal year. Microsoft Cloud revenue surpassed $49 billion, up 26% year-over-year and our commercial RPO increased over 50% to nearly $400 billion with a weighted average duration of only two years. We are seeing increasing demand and diffusion of our AI platform and family of Copilots, which is fueling our investments across both capital and talent.

When it comes to infrastructure, we're building a planet-scale cloud and an AI factory, maximizing tokens per dollar per watt while supporting the sovereignty needs of customers in countries. We are innovating rapidly across the family of Copilots spanning the high value domains of information work, coding, security, science, health and consumer. And as you saw yesterday, we closed a new definitive agreement with OpenAI, marking the next chapter in what is one of the most successful partnerships and investments our industry has ever seen. This is a great milestone for both companies and we continue to benefit mutually from each other's growth across multiple dimensions.

Already, we have roughly 10x-ed our investment. OpenAI has contracted an incremental $250 billion of Azure services. Our rev share, exclusive IP rights and API exclusivity for Azure continue until AGI or through 2030, and we have extended the model and product IP rights through 2032. And we are also energized to innovate and pursue AI advancements with both talent and compute investments that have real-world impact.

With that, let's turn to our momentum across our AI platform and Copilots as well as with agents. We have the most expansive data center fleet for the AI era and we are adding capacity at an unprecedented scale. We will increase our total AI capacity by over 80% this year and roughly double our total data center footprint over the next two years, reflecting the demand signals we see. Just this quarter, we announced the world's most powerful AI data center, Fairwater in Wisconsin, which will go online next year and scale to 2 gigawatts alone.

And we have deployed the world's first large scale cluster of NVIDIA GB300s. We are building a fungible fleet that's being continuously modernized and spans all stages of the AI lifecycle from pre-training to post-training to synthetic data generation and inference and it also goes beyond GenAI workloads to recommendation engines, databases and streaming. We are optimizing this fleet across silicon, systems and software to maximize performance and efficiency. It's this combination of fungibility and continuous optimization that allows us to deliver the best ROI and TCO for us and our customers. For example, during the quarter, we increased the token throughput for GPT-4.1 and GPT-5, two of the most widely used models by over 30% per GPU.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

We also have the most comprehensive digital sovereignty platform. Azure customers in 33 countries are now developing their own cloud and AI capabilities within their borders to meet local data residency requirements. In Germany, for example, OpenAI and SAP will rely on Azure to deliver new AI solutions to the public sector. On top of this infrastructure, we are building Azure AI Foundry to help customers build their own AI apps and agents. We have 80,000 customers, including 80% of the Fortune 500. We offer developers and enterprise access to over 11,000 models more than any other vendor, including as of this quarter, OpenAI's GPT-5 as well as xAI's Grok 4. For example, Ralph Lauren used Foundry to build conversational shopping experience in its app enabling customers to describe what they're looking for and get personalized recommendations. And OpenEvidence used Foundry to create its AI powered clinical assistant, which surfaces relevant medical information to physicians and help streamline charting.

When it comes to our first-party models, we are excited by the performance of our new MAI models for text, voice and image generation, which debuted among the top in the industry leaderboards. And we continue to make great progress with our Phi family of SLMs, which now have been downloaded over 60 million times, up 3x year-over-year. Beyond models and Foundry, we're providing everything developers need to design, customize and manage AI applications and agents at scale. Our new Microsoft Agent Framework helps developers orchestrate multi-agent systems with compliance, observability and deep integration out-of-the-box. For example, KPMG used the Framework to modernize the audit process connecting agents to internal data with enterprise-grade governance and observability. These kinds of real production scale AI deployments are driving Azure's overall growth. And once again, this quarter, Azure took share.

Now, let's turn to applications and agents we ourselves are building on this platform. We now have 900 million monthly active users of our AI features across our products. And our first-party family of Copilots now has surpassed 150 million monthly active users across information work, coding, security, science, health and consumer. When it comes to information work, we continue to innovate with Microsoft 365 Copilot. Copilot is becoming the UI for the agentic AI experience. We have integrated chat and agentic workflows into everyday tools like Outlook, Word, Excel, PowerPoint and Teams. Just nine months since release, tens of millions of users across Microsoft 365 customer base are already using chat. Adoption is accelerating rapidly, growing 50% quarter-over-quarter and we continue to see usage intensity increase. This quarter, we also introduced Agent Mode, which turns single prompts into expert quality Word Documents, Excel Spreadsheets, PowerPoint Presentation and then iterate to deliver the final product, much like Agent Mode in coding tools today. We're thrilled by the early response, including third-party benchmarks that rank it best-in-class.

Beyond individual productivity, Copilot is multi-player with Teams Mode announced this week. You can now invite colleagues into a Copilot conversation. And our collaborative agents like Facilitator and Project Manager prep meeting agendas, take notes, capture decision, kick-off group tasks. We are seeing a growing Copilot agent ecosystem with top ISVs like Adobe, Asana, GRAIL, LexisNexis, SAP, ServiceNow, Snowflake and Workday all building their own agents that connect to Copilot. And customers are also building agents for their mission-critical business processes and workflows using tools like Copilot Studio and integrating them into Copilot.

The overall number of agent users doubled quarter-over-quarter, and just yesterday, we announced App Builder, a new Copilot agent that lets anyone create and deploy task-specific apps and agents in minutes grounded in Microsoft 365 context. All this innovation is driving our momentum. Customers continue to adopt Microsoft 365 Copilot at a faster rate than any other new Microsoft 365 suite, all up more than 90% of the Fortune 500 now use Microsoft 365 Copilot. Accenture, Bristol Myers Squibb, EY Global and the UK's Tax and Payment and Customs Authority all purchased over 15,000 seats this quarter. Lloyds Banking Group has deployed

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

30,000 seats, saving each employee an average of 46 minutes daily. And a large majority of our enterprise customers continue to come back to purchase more seats. Our partner PwC alone added 155,000 seats this quarter and now has over 200,000 deployed across its global operations. In just six months, PwC employees interacted with Microsoft 365 Copilot over 30 million times and they credit this agentic transformation with saving millions of hours in employee productivity.

When it comes to coding, GitHub Copilot is the most popular AI pair programmer now with over 26 million users. For example, tens of thousands of developers at AMD use GitHub Copilot accepting hundreds of thousands of lines of code suggestions each month and crediting it with saving months of development time. All up, GitHub is now home to over 180 million developers, and the platform is growing at the fastest rate in its history, adding a developer every second. 80% of new developers on GitHub start with Copilot within the first week.

Overall, the rise of AI coding agents is driving record usage with over 500 million pull requests merged over the past year. And just yesterday, at GitHub Universe, we introduced Agent HQ. GitHub Copilot and Agent HQ is the organizing layer for all coding agent extending the GitHub primitives like PRs, issues, actions to coding agents from OpenAI, Anthropic, Google, Cognition, xAI as well as OSS and in-house models. GitHub now provides a single mission control to launch, manage and review these agents, each operating from its own branch with built-in controls, observability and governance. We are building a similar system in security with over three dozen agents in Copilot integrated across Entra, Defender, Purview and Intune. For example, with our Phishing Triage Agent and Defender studies show that analysts can be up to 6.5 times more efficient in detecting malicious mails.

In health, Dragon Copilot helps providers automate critical workflows. This quarter alone, we helped document over 17 million patient encounters up nearly 5x year-over-year. More than 650 healthcare organizations have purchased our ambient listening tech to-date, including University of Michigan Health, where over 1,000 physicians are actively using it.

Finally, when it comes to AI consumer experiences, we are excited about all the progress Copilot is making, starting with Windows. Every Windows 11 PC now is an AI PC. Just two weeks ago, we introduced new ways to speak naturally to your computer, including a Copilot wake word, with Vision, Copilot sees what you see on your screen, and you can have a real-time conversation about it. And with Action, it takes real action on your behalf, interacting with both web and desktop apps.

In Edge, we are introducing first of its kind AI features to automate multi-step workflows within the browser and help you pick-up right where you left off. Edge now has taken share for 18 consecutive quarters. In Bing, our overview pages now include embedded conversational capabilities. We took share again in Search. And daily users of our Copilot consumer app increased nearly 50% quarter-over-quarter. Among many updates we made last week is Groups, which turns Copilot for the first time into a shared experience. We also are creating a great consumer subscription offer with Microsoft 365 Premium. It brings together our Office applications and advanced Copilot features with high usage limits, giving individuals the flexibility to bring their own AI to work in a secure way.

Finally, in gaming, Copilot provides a voice-first immersive experience across PC, mobile and our new Xbox Ally. Beyond our family of Copilots and AI platform, we are seeing strong momentum across the portfolio. Cloud migrations are accelerating. In data and analytics, Fabric revenue grew 60%, which is faster than any other data and analytics platform in the industry. We now have 28,000 paid Fabric customers.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

In databases, SQL DB, hyperscale revenue was up nearly 75%, 50% in Cosmos DB. In business applications, Dynamics 365 gained share. In security, our end-to-end stack is now informed by 100 trillion daily signals, 1 billion monthly active users of Entra, 16 billion Copilot interactions audited by Purview, up 72% quarter-over-quarter, 40,000 sentinel customers and we took share across all categories we serve in security.

In LinkedIn, nearly 1.3 billion members. And finally, in gaming, we expanded our reach across every endpoint focused on our high margin content and services. We launched critically acclaimed games like Keeper, Ninja Gaiden 4 and Outer Worlds 2, reaching 155 million monthly active users in Minecraft, an all-time high and set new record for overall content and services revenue for the quarter. We also saw a great response to Xbox Ally launch two weeks ago and set new records for players on PC.

In closing, our planet-scale Cloud and AI factory together with Copilots across high-value domains is driving broad diffusion and real-world impact, and we continue to increase our investments in AI across both capital and talent to meet the massive opportunity ahead.

With that, let me turn it over to Amy to walk through our financial results and outlook, and I look forward to rejoining for your questions.

Amy Hood {BIO 18040963 <GO>}

Thank you, Satya, and good afternoon, everyone. First, as you heard from Satya, we were pleased to announce the next phase of our partnership with OpenAI yesterday. They continue to choose Microsoft to power their workloads, and together, we remain committed to driving innovation that meets real-world needs. Our Q1 results were not impacted by the deal signed this week.

Now, onto the quarter. We delivered a strong start to our fiscal year, exceeding expectations across revenue, operating income and earnings per share. We also saw continued share gains across many of our businesses, demonstrating our leadership position in key markets. This quarter, revenue was $77.7 billion, up 18% and 17% in constant currency. Gross margin dollars increased 18% and 16% in constant currency, while operating income increased 24% and 22% in constant currency. And earnings per share was $4.13, an increase of 23% and 21% in constant currency when adjusted for the impact of our investments in OpenAI.

FX impact was roughly in line with guidance. Company gross margin percentage was 69%, down slightly year-over-year, driven by investments in AI, including the impact of scaling our AI infrastructure and the growing usage of our AI product features. This was partially offset by ongoing efficiency gains, particularly in Azure and M365 Commercial cloud.

Operating expenses increased 5% and 4% in constant currency, driven by investments in Cloud and AI engineering, including compute capacity and AI talent to support product development across the portfolio. Operating margins increased year-over-year to 49% and were ahead of expectations with stronger than anticipated results in high-margin businesses this quarter.

When adjusted for the impact from our investments in OpenAI, other income and expense was $401 million as interest income more than offset interest expense, which includes the interest payments related to datacenter finance leases. Capital expenditures were $34.9 billion, driven by growing demand for our cloud and AI offerings. This quarter, roughly half of our spend was on short-lived assets, primarily GPUs and CPUs to support increasing Azure platform demand, growing first-party apps and AI solutions, accelerating R&D by our product teams as well as continued replacement for end-of-life server and networking equipment.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

The remaining spend was for long lived assets that will support monetization for the next 15 years and beyond, including $11.1 billion of finance leases that are primarily for large datacenter sites. And cash paid for PP&E was $19.4 billion. As a reminder, the difference between total CapEx and cash paid for PP&E is primarily due to finance leases as well as the normal timing of goods received but not yet paid.

Cash flow from operations was $45.1 billion, up 32%, driven by strong cloud billings and collections, partially offset by higher supplier payments. And free cash flow increased 33% to $25.7 billion with minimal impact from a sequential increase in CapEx given the higher mix of finance leases. And finally, we returned $10.7 billion to shareholders through dividends and share repurchases.

Now, to our Commercial results. Commercial bookings increased 112% and 111% in constant currency and were significantly ahead of expectations, driven by Azure commitments from OpenAI as well as continued growth in the number of $100 million plus contracts for both Azure and M365. These results do not include any impact from the incremental $250 billion Azure commitments from OpenAI announced yesterday. Commercial remaining performance obligation increased to $392 billion and was up 51% year-over-year. The balance has nearly doubled over the past two years. And even with this growth, our weighted average duration has been relatively stable at approximately two years.

Microsoft Cloud revenue was $49.1 billion ahead of expectations and grew 26% and 25% in constant currency. Microsoft Cloud gross margin percentage was slightly better than expected at 68% and down year-over-year due to the investments in AI that were partially offset by ongoing efficiency gains as noted earlier.

Now, to segment results. Revenue from productivity and business processes was $33 billion and grew 17% and 14% in constant currency. M365 Commercial Cloud revenue increased 17% and 15% in constant currency with one point of benefit from in period revenue recognition. Year-over-year growth was driven by both ARPU and seats, with ARPU growth again led by E5 and M365 Copilot. Paid M365 Commercial seats grew 6% year-over-year with installed base expansion across all customer segments, though primarily in our small and medium businesses and frontline worker offerings.

M365 Commercial products revenue increased 17% and 14% in constant currency ahead of expectations due to higher-than-expected Office 2024 transactional purchasing. M365 Consumer Cloud revenue increased 26% and 25% in constant currency, again driven by ARPU growth. M365 consumer subscriptions grew 7% to over $90 million.

LinkedIn revenue increased 10% and 9% in constant currency driven by marketing solutions. The Talent Solutions business was impacted by continued weakness in the hiring market. Dynamics 365 revenue increased 18% and 16% in constant currency with continued growth across all workloads. Segment gross margin dollars increased 19% and 16% in constant currency and gross margin percentage increased, driven by efficiency gains in M365 Commercial Cloud that were partially offset by investments in AI, including the impact of growing usage in M365 Copilot chat.

Operating expenses increased 6% and 5% in constant currency and operating income increased 24% and 20% in constant currency. Operating margins increased 3 points year-over-year to 62%, driven by the higher gross margin noted earlier as well as improved operating leverage.

Next, the Intelligent Cloud segment. Revenue was $30.9 billion and grew 28% and 27% in constant currency. In Azure and other cloud services, where we continue to see accelerating demand, revenue grew 40% and 39% in constant currency. Results were ahead of expectations

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

driven by better-than-expected growth in our core infrastructure business, primarily from our largest customers. Azure AI services revenue was generally in line with expectations, and this quarter demand again exceeded supply across workloads even as we brought more capacity online.

In our on-premise server business, revenue increased 1% and was relatively unchanged in constant currency. Results were ahead of expectations driven by transactional purchasing of Windows Server 2025. Segment gross margin dollars increased 20% and 19% in constant currency and gross margin percentage decreased year-over-year, driven by investments in AI that were partially offset by efficiency gains in Azure. Operating expenses increased 4% and operating income grew 27%. Operating margins were 43%, down only slightly year-over-year as increased investments in AI were mostly offset by improved operating leverage.

Now, to More Personal Computing. Revenue was $13.8 billion and grew 4%. Windows OEM and Devices revenue increased 6% year-over-year, significantly ahead of expectations, driven by strong demand ahead of Windows 10 end of support as well as a benefit from inventory levels that remain elevated. Search and news advertising revenue ex-TAC increased 16% and 15% in constant currency, driven by growth in volume as well as a continued benefit from third-party partnerships that was better than expected.

And in gaming, revenue decreased 2% and 3% in constant currency. Against a strong prior year comparable, Xbox content and services revenue increased 1% and was relatively unchanged in constant currency, driven by better-than-expected performance from third-party content. Segment gross margin dollars increased 11% and 10% in constant currency and gross margin percentage increased year-over-year, driven by sales mix shift to higher-margin businesses. Operating expenses increased 4% and 3% in constant currency and operating income increased 18% and 16% in constant currency. Operating margins increased 3 points year-over-year to 30%, driven by the higher gross margin noted earlier.

Now, moving to our Q2 outlook, which unless specifically noted otherwise is on a US dollar basis. Based on current rates, we expect FX to increase total revenue growth by 2 points. Within the segments, we expect FX to increase revenue growth by 2 points in productivity in Business Processes and Intelligent Cloud and 1 point in More Personal Computing. We expect FX to increase COGS and operating expense growth by 1 point. Starting with the total company, we expect revenue of USD79.5 billion to USD80.6 billion or growth of 14% to 16%. We expect COGS of USD26.35 billion to USD26.55 billion or growth of 21% to 22% and operating expense of USD17.3 billion to USD17.4 billion or growth of 7% to 8%. Operating margins should be relatively flat year-over-year and down sequentially in line with historic seasonality.

Now other income and expense. The combination of OpenAI's conversion to a public benefit Corp and the ongoing nature of our partnership will result in increased volatility. Therefore, going forward, we will provide our outlook excluding any impact from our investments in OpenAI. On that basis, in Q2, other income and expense is estimated to be roughly $100 million as interest income will more than offset interest expense. And we expect our Q2 effective tax-rate to be approximately 19%.

Next, capital expenditures. With accelerating demand and a growing RPO balance, we're increasing our spend on GPUs and CPUs. Therefore, total spend will increase sequentially and we now expect the FY'26 growth rate to be higher than FY'25. As a reminder, there can be quarterly spend variability from cloud infrastructure build outs and the timing of delivery of finance leases.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

Next, our Commercial business. In Commercial bookings, we expect healthy growth in the core business on a low expiry base when adjusted for the OpenAI contracts in the prior year. And we expect Commercial bookings will be positively impacted by the significant OpenAI commitments announced yesterday. As a reminder, larger long-term Azure contracts, which are more unpredictable in their timing drive increased quarterly volatility in our bookings growth rate. Microsoft Cloud gross margin percentage be roughly 66%, down year-over-year, driven by the continued investments in AI as well as the mix-shift to Azure.

Now, to segment guidance. In productivity and business processes, we expect revenue of USD33.3 billion to USD33.6 billion or growth of 13% to 14%. In M365 Commercial Cloud, we expect revenue growth to be between 13% and 14% in constant currency with business trends that remain relatively stable quarter-over-quarter. ARPU growth will again be driven by E5 and M365 Copilot. M365 Commercial products revenue growth should be in the low to mid-single-digits. As a reminder, M365 Commercial products, includes components that can be variable due to the in-period revenue recognition dynamics. M365 consumer cloud revenue growth should be in the mid 20s driven by growth in ARPU.

For LinkedIn, we expect revenue growth of approximately 10%. And in Dynamics 365, we expect revenue growth to be in the mid to high teens with continued growth across all workloads. For Intelligent Cloud, we expect revenue of USD32.25 billion to USD32.55 billion or growth of 26% to 27%. In Azure, we expect Q2 revenue growth of approximately 37% in constant currency as demand remains significantly ahead of the capacity we have available. And while we're accelerating the amount of capacity we're bringing online, we will continue to balance Azure revenue growth with the growing needs across our first-party apps and AI solutions, our own R&D efforts and the end-of-life server replacements. Therefore, we now expect to be capacity-constrained through at least the end of our fiscal year. As a reminder, there can be quarterly variability in the year-on-year growth rates depending on the timing of capacity delivery and when it comes online, as well as from in-period revenue recognition depending on the mix of contracts.

In our on-premises server business, we expect revenue to decline in the low to mid-single-digits with ongoing customer shift to cloud offerings. In More Personal Computing, we expect revenue to be in the USD13.95 billion to USD14.45 billion. Windows OEM and Devices revenue should decline in the mid-single-digits. We expect continued momentum from Windows 10 end-of-support, although growth rates will be impacted by elevated inventory levels at the end of Q1 that we expect to come down through the quarter. Therefore, Windows OEM revenue should decline low to mid-single-digits. The range of potential outcomes remains wider than normal. Devices revenue should decline year-over-year.

Search and news advertising ex-TAC revenue growth should be in the low double-digits, down sequentially as growth rates normalize following the benefit from third-party partnerships noted earlier. Growth will continue to be driven by volume and revenue per search across Edge and Bing. And in Xbox content and services, we expect revenue to decline the low to mid-single-digits against a prior year comparable that benefited from strong first-party performance, partially offset by growth in subscriptions and hardware revenue should decline year-over-year.

And in closing, demand signals across bookings, RPO and product usage are accelerating faster than we expected. We're investing in infrastructure, AI talent and product innovation to capture that momentum and expand our leadership position. And we remain focused on delivering real value to our customers that results in durable revenue growth for the long-term.

With that, let's go to Q&A, Jonathan.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

Jonathan Neilson

Thanks, Amy. We'll now move over to Q&A. Out of respect for others on the call, we request that participants please only ask one question. Operator, can you please repeat your instructions?

Questions And Answers

Operator

(Operator Instructions). And our first question comes from the line of Keith Weiss with Morgan Stanley. Please proceed.

Q - Keith Weiss {BIO 6993337 <GO>}

Excellent. Thank you guys for taking the question and congratulations on another outstanding quarter. And if I'm looking at Microsoft, this is two quarters in a row, we're really seeing results that are well ahead of anybody's expectations when we were thinking about this company a year-ago or five years ago. 111% Commercial bookings growth was not on anybody's bingo card, if you will. Yet the stock is underperforming the broader market. And the question I have is kind of getting at the zeitgeist that I think is weighing on the stock. And it's -- is something about to change? And I think AGI is kind of a nomenclature or shorthand for that and it's something that's still included in your guys' OpenAI agreement.

So Satya, when we think about AGI or we think about how application and computing architectures are changing, is there anything that you see on the horizon, whether it's AGI or something else that could potentially change what appears to be a really strong positioning for Microsoft in the marketplace today where that strength will perhaps weaken on a go-forward basis. Is there anything that you're worrying about in that evolution and particularly the evolution of these generative AI models?

A - Satya Nadella {BIO 3224315 <GO>}

Yeah, thank you, Keith, for the question. So here's how I'd say that I think there are two parts. We feel very, very good about even this, I'd say, the new agreement that we now have with OpenAI because I think even it just creates more certainty to all the IP relationship we have as it relates to even this definition of AGI. But beyond that, I think your question touches on something that's pretty important, which is, how are these AI systems going to truly be deployed in the real-world and make a real difference and make a return for both the customers, who are deploying them, and then obviously the providers of these systems.

And I think the best way to characterize the situation is that even as the intelligence capability increases, let's even say exponentially by model version over model version. The problem is, it's always going to still be jagged, right? I think the term people use is the jagged intelligence. Even it or spiky intelligence, right? So you may even have a capability that's fantastic at a particular task, but it may not uniformly grow. So, what is required is, in fact, these systems, whether it is GitHub Agent HQ or the M365 Copilot system, don't think of this as a product, think of it as a system that in some sense smooths out those jagged edges and really helps the capability.

I mean, just to give you a flavor for it, right? So if I am in M365 Copilot, I can generate an Excel spreadsheet. The good news is now an Excel spreadsheet does understand Office JS, has the formulas in it. And it feels like, it is a great spreadsheet created by a good model. The more interesting thing is I can go into Agent Mode in Excel and iterate on that model. And yet it will stay

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

on rails. It won't go off rails. It will be able to do the iteration. Then I can even give it to the analyst agent and then it will even make sense of it like a data analyst would of our Excel model.

The reason I say all of that is because that's the type of construction that will be needed. Even when the model is magical, all powerful, I think we will be in this jagged intelligence phase for a long time. So one of the fundamental things that these -- whether it's GitHub, whether it's security, whether it's M365, the three main domains we're in, we feel very, very good about building these as organizing layers for agents to help customers. And by the way, that's the same thing that we want to put into Foundry for our third-party customers. So that's kind of how people will build these multi-agent systems.

So I feel actually pretty good about both the progress in AI. I don't think AGI as defined at least by us in our contract is ever going to be achieved anytime soon. But I do believe we can drive a lot of value for customers with advances in AI models by building these systems. So it's kind of the real question that needs to be well understood, and I feel very, very confident about our ability to make progress.

Q - Keith Weiss {BIO 6993337 <GO>}

Excellent. That's super-helpful.

A - Jonathan Neilson

Thanks, Keith. Operator, next question, please.

Operator

The next question comes from the line of Brent Thill with Jefferies. Please proceed.

Q - Brent Thill {BIO 1556691 <GO>}

Thanks. Amy, on the bookings blowout, I guess many are somewhat concerned about concentration risk. And I think you noted a number of $100 million contracts, not to go into a lot of detail, but can you just give us a sense of what you're seeing in that 51% RPO and 110% plus bookings growth that gives you confidence about what you're seeing in terms of the breadth and you know extent of some of these deals on a global basis? Thanks.

A - Amy Hood {BIO 18040963 <GO>}

Yeah. Thanks, Brent. A couple of things to maybe take a step back, on RPO, with a nearly $400 billion balance, we've been trying to help people understand sort of how to think about really the breadth of that. It covers numerous products, it covers customers of all sizes, that's been a balance that we've been growing, obviously at a good clip. But what people need to realize is it sits across multiple products because of the things Satya is talking about around creating systems and where we're investing. And if you're going to have that type of balance, and then more importantly, have the weighted average duration be two years, it means that most of that is being consumed in relatively short order. People are not consuming, and I say this broadly, unless there is value. And I think this is why we keep coming back to are we creating real-world value in our AI platforms, in our AI solutions and apps and systems.

And so, I think the sort of the way to think about RPO is it's been building across a number of customers. We're thrilled to have OpenAI be a piece of that. We're learning a ton and building leading systems because of it that are being used at-scale that benefits every other customer. And so, it's why we've tried to give a little bit more color to that RPO balance because I do

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

understand that there have been a lot of concerns or questions about is it long dated, is it coming over a long period of time, and hopefully this is helpful for people to realize that these are contracts being signed by customers who intend to use it in relatively short order. And at that type of scale, I think that's a pretty remarkable execution.

Q - Brent Thill {BIO 1556691 <GO>}

Thank you.

A - Jonathan Neilson

Thanks, Brent. Operator, next question, please.

Operator

The next question comes from the line of Mark Moerdler with Bernstein Research. Please proceed.

Q - Mark Moerdler {BIO 16855032 <GO>}

Thank you very much for taking my question and congratulations on the quarter. It's pretty amazing what you guys are doing. Satya and Amy, I'd like to ask you the number one question I receive whether from investors or at AI conferences I attend. How much confidence do you have that the software and even the consumer Internet business can monetize all the investments we're seeing globally or frankly, are we in a bubble? In fact, Amy, what would be the factors you'd be watching for to assure that you're not overbuilding for current demand and that demand will sustain? Thank you.

A - Amy Hood {BIO 18040963 <GO>}

Maybe I'll start, Satya, and then you could add. Let me talk a little bit about maybe connecting a couple of the dots because with $400 billion of RPO, that's sort of short dated as we talked about, our needs to continue to build out the infrastructure is very high. And that's for booked business today. That is not any new booked business we started trying to accomplish on October 1st, right?

And so, the way to think about that, and you saw it this quarter in particular, and as we talked about '26, the remainder. Number one, we're pivoting toward increasingly, we talked about these short-lived assets, both GPUs and CPUs. Again, we talk about all these workloads are burning both in terms of app building. Now when that happens, short-lived assets generally are done to match sort of the duration of the contracts or the duration of your expectation of those contracts. And so, I sometimes think when people think about risk, they're not realizing that most of the lifetimes of these and the lifetimes of the contracts are very similar.

And so, when you think about having revenue and the bookings and coming on the balance sheet and the depreciation of short-lived assets, they're actually quite matched, Mark. And as you know, we have spent the past few years not actually being short GPUs and CPUs per se, we were short the space or the power is the language we use to put them in. So we spent a lot of time building out that infrastructure. Now we're continuing to do that also using leases. Those are very long-lived assets as we've talked about, 15 years to 20 years. And over that period of time, do I have confidence that we'll need to use all of that, it is very high.

And so, when I think about sort of balancing those things, seeing the pivot to GPU, CPU, short-lived, seeing the pivot in terms of how those are being utilized, we are, and I said this now, we've

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

been short now for many quarters, I thought we were going to catch-up. We are not. Demand is increasing. It is not increasing in just one place. It is increasing across many places. We're seeing usage increases in products. We are seeing new products launch that are getting increasing usage and increasing usage very quickly. You know, when people see real value, they actually commit real usage. And I sometimes think this is where this cycle needs to be thought through completely is that when you see the kind of demand signals and we know we're behind, we do need to spend. But we're spending with a different amount of confidence in usage patterns and in bookings. And I feel very good about that.

I have said, we are now likely to be short capacity to serve the most important things we need to do, which is Azure, our first-party applications. We need to invest in product R&D and we're doing end-of-life replacements in the fleet. So we're going to spend to make sure that happens. It's about modernization, it's about high-quality, it's about service delivery and it's about meeting demand. And so, I feel good about doing that and I feel good that we've been able to do it so efficiently and with a growing book of business behind it.

A - Satya Nadella {BIO 3224315 <GO>}

Yeah. The only thing I'd add to what Amy captured was, if you sort of look out, there are two things that matter, I think, and that are critical in terms of how we think about our allocation of capital, also our R&D. One is, how efficient is our planet-scale token factory, right? I mean that's at the end of the day what you have to do. And in order to do that, you have to start with building out a very fungible fleet. It's not like we are building one data center in one region in the world that's mega-scale. We are building it out across the globe for inference, for pre-training, for post-training, for RL, for datacenters and what have you. So therefore, that's the fungibility is super important.

The second thing that we're also doing is continually modernizing the fleet. It's not like we buy one version of, say, NVIDIA and load-up for all the gigawatts we have. Each year you buy, you write the Moore's Law, you continuously modernize and depreciate it. And that means you also use software to grow efficiency. I talked about, I think, 30% improvement on both serving up GPT-4.1 and GPT-5, right? That's software. That's sort of -- and by the way, it's helpful on A100s. It's helpful on GB200s and it will be helpful on GB300s. That's the beauty of having the efficiency of the fleet. So keep improving utilization, keep improving the efficiency. So that's what you do in the token factory.

The other aspect, which Amy spoke to is, we have some of the best agent systems that matter in the high-value domains, right? In information work, that's the copilot system. Coding, I mean, I should also say one of the things I like about Copilot is, I mean Copilot ARPUs compared to M365 ARPUs, right, it's expansive. The same thing that happened between server and cloud, like we used to always say, well, is it zero-sum? It turned out that the cloud was so much more expansive to the server market. The same thing is happening in AI because first you could say, hey, our ARPUs are too low when it comes to M365, or you could say we have the opportunity with AI to be much more expansive.

Same thing with tools, right? I mean, tooling -- tools business was not like a leading business, whereas coding business is going to be one of the most expansive AI systems. And so, we feel very good about being in that category. Same thing with security, same thing with health. So we have -- and in consumer, one of the things is it's not just about ads, it's ads plus subscriptions that also opens up opportunity for us. So when I look at the entirety of these high-value agent systems and when we look at the efficiency and fungibility of our fleet, that's what gives us the confidence to invest both the capital and the R&D talent to go after this opportunity.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

Q - Mark Moerdler {BIO 16855032 <GO>}

That was pretty amazing. I really appreciate all the detail.

A - Jonathan Neilson

Thanks, Mark. Operator, next question please.

Operator

The next question comes from the line of Karl Keirstead with UBS. Please proceed.

Q - Karl Keirstead {BIO 1542979 <GO>}

Okay. Thank you. This one is for Amy. Amy, I certainly don't want to take you down a too complex and accounting path with this question. But the investment in OpenAI that sits in other income at $4.1 billion. It is so large that I think the audience could -- listening-in could benefit from a little bit more color about what that is. It feels like it's so much larger than you were running through other income in prior quarters that it mustn't just be your share of the OpenAI losses. So could you just describe that and what we can expect in subsequent quarters and whether this signals any kind of accounting change? Thanks so much.

A - Amy Hood {BIO 18040963 <GO>}

The Q1 number was not impacted at all by the new agreement that was put in place, let me first say that. Secondly, that increased loss was all due to our percentage of losses at OpenAI due to the equity method. So just to be very clear. So there is not anything there that is not the increased losses from OpenAI.

Q - Karl Keirstead {BIO 1542979 <GO>}

Okay. Understood. Thank you.

A - Amy Hood {BIO 18040963 <GO>}

Thanks, Karl.

A - Jonathan Neilson

Thanks, Karl. Operator, next question please.

Operator

The next question comes from the line of Mark Murphy with J.P. Morgan. Please proceed.

Q - Mark Murphy {BIO 18840231 <GO>}

Thank you so much. So we seem to be entering into a new era where the contractual commitments from a small number of AI natives are just incredibly large, not only in absolute terms, but sometimes relative to the size of the companies themselves. For instance, contracts worth hundreds of billions of dollars that are 20 times their current revenue scale. Philosophically, how do you evaluate the ability of those companies to follow through on these commitments? And how do you think about placing guardrails on customer concentration for any single entity?

A - Satya Nadella {BIO 3224315 <GO>}

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

Yeah. Maybe I'll start and then Amy you can add. I mean it goes back a little bit, Mark, to what I said about building first, the asset itself such that it's most fungible. And then to recognize the strength of even sort of our portfolio. We have a third-party business, we have a first-party business, we have third-party also spread between enterprise, digital natives. I've always felt that we need a balance there because it may start with digital natives. They're always going to be the early adopters. You always have the hit app of the generation. And then, essentially then it spreads throughout the enterprise adoption cycle is just starting. And so therefore, having the -- over the arc of time, I think that third-party balance of customers will only increase.

But it's great to have the hit first-party apps in the beginning because you can build scale that then if it's a fungible and that's where the key is. You don't want to build for a digital native in as if you were just doing hosting for them. You want to build, that's where I think some of the decision-making of ours is probably getting better understood. What do we say yes to, what do we say no to. I think there was a lot of confusion. Hopefully by now anyone who switched on would figure this out. And so that's I think, one thing we're doing on the third-party.

But the one first-party is probably where a lot of our leverage comes and it's not even about one hit app on our first party even. Our portfolio of stuff, which I just walked through in the earlier answer gives us again the confidence that between that mix, we will be able to use our fleet to the maximum. And remember, these assets, especially the datacenters and so on, are long assets, right? There will be many refresh cycles for any one of these when it comes to the gear. So I feel that once you think about all those dimensions, the concentration risk gets mitigated by being thoughtful about how you really ensure the build is for the broad customer base.

A - Amy Hood {BIO 18040963 <GO>}

And maybe just to help with another angle, that I think Satya has helped a lot is that when you think about concentration risk for delivering to any customer, you have to remember that because we're talking about this very large flexible fleet that can be used for anyone and for any purpose, 1P, 3P and including our Commercial Cloud, by the way, which I should be quite clear on, is pretty flexible in every regard. You have to remember that the CPU and GPU and the storage gear doesn't come into play until the contracts start happening. And so, you're right, some of these large contracts have delivery dates over time. So you get a lot of lead time in being able to say, oh, what's the status? So, I think we're pretty thoughtful around what's always gone in our RPO balance and been considerate of that. There's always been that taken into account when we publish that bookings number and publish the RPO balance.

Q - Mark Murphy {BIO 18840231 <GO>}

Thank you very much.

A - Jonathan Neilson

Thanks, Mark. Operator, next question, please.

Operator

The next question comes from the line of Brad Zelnick with Deutsche Bank. Please proceed.

Q - Brad Zelnick {BIO 16211883 <GO>}

Great. Thanks so much for taking the question, and I'll echo my congrats on an amazing start to the year. Amy, is there any way to quantify or frame the revenue impact of Azure being short on

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

capacity? And while I appreciate the constraints you face are broad across the industry, is there risk of workloads going elsewhere? And how do you mitigate that?

A - Amy Hood {BIO 18040963 <GO>}

Yeah, Brad, It's a great question. It's always hard to quantify precisely what would have been the revenue impact in quarter. But I would offer a way to think about it is Azure probably does bear most of the revenue impact, because when you think about real priorities that you have to fill first, it's obviously the increasing usage and adoption and sales we've seen of M365 Copilot and the usage of Copilot Chat, which we've seen very different patterns, which we're encouraged by. It's the adoption of security features. It's the GitHub momentum. And so, when you're thinking about it, that is where and it is a priority for us to allocate resourcing there first.

And so, you're right to ask how do I think about that? We've worked very hard to try to mitigate it as best we can, but we have been short in Azure, and we've been clear on it. And I would say the other two priorities that I haven't mentioned maybe as much before is also just making sure our product teams and the AI talent that we've been able to hire into the company really over the past year and a half have access also to significant capacity because we're seeing it make the product better in a loop that is adding great benefit today into products people are using today for real-world work. And so, we are making that a priority to make sure our research teams have that as well as our product engineering teams. And yes, it does impact Azure directly. That is the place where you see that prioritization. But I think it's probably hard for me to give an exact number, but it is safe to say that the number could be higher.

Q - Brad Zelnick {BIO 16211883 <GO>}

Great. Thank you.

A - Jonathan Neilson

Thanks, Brad. Operator, we have time for one last question.

Operator

The last question will come from the line of Kash Rangan with Goldman Sachs. Please proceed.

Q - Kash Rangan {BIO 22095432 <GO>}

Thank you very much. Amy, I just wanted to congratulate you. I think you said before that it is possible to accelerate Azure growth while getting efficient margins and you've done it. Congrats on that. I have one for you, Satya. With respect to the elephant in the room, following -- just being a little more direct following up on Keith Weiss' question, there's talk that another hyperscaler came in and took away the business that was rightfully Microsoft's. I'm sure that there is a different point-of-view here. I'm wondering if you could offer some perspective on your criteria to, is it about a certain volume of business that you wish to execute on the Microsoft paper?

Or is something broader than that I don't think maybe people fully appreciate the terminal value that Microsoft will have on its balance sheet at the end of these contracts, which I think is probably being underestimated as you have a full stack and you've got the multiple vectors to monetize, be it databases, Foundry and to your point that you are a platform company, not just a hyperscaler, maybe that's what it is all about or maybe there's another story about you letting the other hyperscaler company come in from nowhere and claiming a big piece of that four to five-year puzzle. Thank you so much once again. Really appreciate it. And congratulations.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

A - Satya Nadella {BIO 3224315 <GO>}

Well, thanks, Kash. I mean, for us, again, it just always goes back to, I think the core principle, which is build a fleet that is fungible across the planet and works for third-party and first-party and research. So that's essentially what we have done. And so, when some demand comes in shapes that don't fit that goal where it's too concentrated, not just by customer, by location, by type of skewing, right? I think Amy mentioned some very key things. When you think about the margin profile of a hyperscaler, you've got to remember this the AI accelerator piece, but there's compute, there's storage.

And so, if all of the demand just comes for just one meter, that's really not a long-term business we want to be in. That's even from a third-party. We have to balance it with all of our first-party stuff because that's after all a different margin stack for us. And then we have to fund our own R&D and model capability because in the long run, that's what's going to differentiate us. And so, I look at all of those, we sort of use all of that to make sure we are saying yes to all the demand that we want. We say no to some of the demand that may be something that we could serve, but it's not in our long-term interest. And so, that's sort of the decision-making we have done and we feel very, very good about the decisions. In some sense, I feel even each time we say no to the day after I feel better.

A - Amy Hood {BIO 18040963 <GO>}

And Kash, maybe just a -- yeah, and just Kash, I think this is our last call with you, I just want to say thanks and congratulations. It's been a privilege to work with you and best of luck.

A - Satya Nadella {BIO 3224315 <GO>}

Let me add to that. Best of luck, Kash.

Q - Kash Rangan {BIO 22095432 <GO>}

Thanks. Thank you so much. Very kind of you.

A - Jonathan Neilson

Thanks, Kash. That wraps up the Q&A portion of today's earnings call. Thank you for joining us today, and we look forward to speaking with all of you soon.

A - Satya Nadella {BIO 3224315 <GO>}

Thank you, all.

Operator

Thank you. This concludes today's conference. You may disconnect your lines at this time. And thank you for your participation.

This transcript may not be 100 percent accurate and may contain misspellings and other inaccuracies. This transcript is provided "as is", without express or implied warranties of any kind. Bloomberg retains all rights to this transcript and provides it solely for your personal, non-commercial use. Bloomberg, its suppliers and third-party agents shall have no liability for errors in this transcript or for lost profits, losses, or direct, indirect, incidental, consequential, special or punitive damages in connection with the furnishing, performance or use of such transcript. Neither the information nor

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

any opinion expressed in this transcript constitutes a solicitation of the purchase or sale of securities or commodities. Any opinion expressed in the transcript does not necessarily reflect the views of Bloomberg LP. © COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.

© COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.